RUNPOD PROOF OF USE

Team: Taygood
Project: Kachingg: credit-control assistant for the Orbyte hub
Repository: https://github.com/YapHongSanHansen/Xmum (public)

Statement: RunPod was not used.

Our team did not use RunPod. Kachingg's only AI step reads one short message and one slip image per customer reply, so it calls a hosted model (ILMU Mini v3.3 via Grafilab's Models API) rather than running a GPU model. The rules engine needs no GPU.

- **Why not self-host now:** the step is small and infrequent, and a hosted API gave us image input and structured output on day one.
- **Where RunPod would fit:** in a real deployment, Malaysian data-residency preferences (PDPA cross-border transfer) could justify an open-weight vision model on a GPU endpoint such as RunPod Serverless. The /api/extract route is the single place that would change.

Team leader: Hansen Yap Hong San

Signature: _____

Date: 25 October 2026